

Main Concepts

Terraform

Agenda

- Terraform CLI
- Providers & Provisioners
- Resources
- Variables & Outputs
- Organize your code
- Terraform State

Terraform CLI

Terraform

Terraform CLI

- Terraform CLI (Command Line Interface) is a program written in Golang
- Available for all most used platforms (Windows, Linux, Mac, ARM, ...)
- Allow the execution of manual commands but more important to be run inside automation processes
- Use Terraform code (files with .tf extension) as input for each command
- Some available commands
 - terraform init
 - terraform validate
 - terraform fmt
 - terraform plan
 - terraform apply
 - terraform destroy

Providers & Provisioners

Terraform

Terraform Providers

- Terraform allows you to build infra in different clouds but is not cloud agnostic
- Terraform embraces different cloud providers to allow building infra using the same language (HCL) and same principles
- But when you have Terraform code to build an environment in Azure, if you need to create in AWS too, you need to create new code from scratch
- What allows Terraform to communicate with all cloud providers, are the Terraform Providers

Terraform Providers

- You can have Terraform Providers for cloud providers, SaaS applications, on-prem virtualization, etc.
- Every solution that you use that gives you a CLI or a well-defined API may have a Terraform Provider
- If you need a customized provider, you can build your own using Golang
- Examples
 - Cloud providers: Azure, AWS, GCP
 - DevOps Platforms: GitHub, Gitlab, Azure DevOps
 - Monitoring: NewRelic
 - VPNs: Twingate, Tailscale

Terraform Provisioners

- A Terraform Provider can be seen as a library that you use on your code and allow you to write code to use it
- But then, you need to translate that code to something that your final provider knows
- For example, if you use Azure provider to write the code to create a VM, you need to translate from that code to Azure instructions
- That's the role of the provisioners!
- Using Azure example, the provisioner receives the Terraform code and translate for **az CLI** commands to reach your Azure Resource Manager
- These provisioners need to be installed by **terraform init** code and must be always the first command to be run

How to define providers?

- You can add as many providers you want to your Terraform code
- For each provider you should always specify a version and specific properties needed by the provider
- The version should be defined on a terraform config block to centralize all needed versions configuration

```
terraform {  
  required_version = ">= 1.4.0"  
  required_providers {  
    azurerm = "3.55.0"  
  }  
}  
  
provider "azurerm" {  
  subscription_id = "subscription-id"  
  client_id = "principal-used-for-access"  
  client_secret = "password-of-principal"  
  tenant_id = "tenant-id"  
  alias = "arm-1"  
}
```

How to set version numbers?

- On Terraform block you should specify:
 - Terraform CLI version to be used with your code
 - Version for each provider your using
- When you define a Terraform CLI version, you may have a CLI version compliant with defined on the code
- How can version number be defined:
 - '=': Allows only one exact version number
 - '!=': Excludes an exact version number
 - '<', '<=', '>', '>=': Compares the version and allows any when comparison is true. '>= 3.0.0' allows '3.0.0', '3.0.1', '3.1.0' and '4.0.0'
 - '~>': Allows only the rightmost version component to increment
 - '~> 1.0.4': Allows versions '1.0.5' and '1.0.10' but not '1.1.0'
 - '~> 1.1': Allows versions '1.2' and '1.10' but not '2.0'

Resources

Terraform

Resources

- Resources are the building blocks of Terraform code
- When you author Terraform code you're defining resources and the relationship between them
- Each resource represents an element on the target provider
- Inside each resource you define unique properties that make sense on the context of that resource

Resources

- **resource** = Top level keyword
- **type** = Type of resource. Ex: `azurerm_virtual_machine`.
- **name** = Arbitrary name to refer to this resource. Used internally by terraform. This field cannot be a variable.

```
resource "type" "name" {  
    parameter = "foo"  
    parameter2 = "bar"  
    list = ["one", "two", "three"]  
}
```

Resources: Example

- Resource to create AWS VM
- **aws_instance** = Type of resource available on AWS provider
- **ex** = name to the resource
- **ami** and **instance_type** = properties of the AWS VM

```
resource "aws_instance" "ex"{  
    ami = "ami-c58c1dd3"  
    instance_type = "t2.micro"  
}
```

Resources: Dependency Mapping

- By default, Terraform tries to create all resources in parallel
- But sometimes that is not possible because you have dependency between resources
- Dependency mapping is a out-of-the-box feature of Terraform that uses the code you created to define the sequence to create the resources
- Based on your code Terraform creates a dependency graph between all resources and start to create resources in that sequence

Resources: Dependency Mapping

- In some cases, mostly due to provider implementation, you may need to add an explicit dependency between resources
- Terraform detects circular references too and gives an error (without changing any resource) when that happens

```
resource "azurerm_resource_group" "hashitraining" {
  name      = "${var.prefix}-vault-workshop"
  location  = var.location
}

resource "azurerm_virtual_network" "vnet" {
  name            = "${var.prefix}-vnet"
  location        = azurerm_resource_group.hashitraining.location
  address_space   = [var.address_space]
  resource_group_name = azurerm_resource_group.hashitraining.name
}
```


Resources: How to reference?

- The reference from one resource to the other, follows this pattern

resource_type.resource_name.property_name

```
resource "azurerm_resource_group" "hashitraining" {  
  name      = "${var.prefix}-vault-workshop"  
  location = var.location  
}  
  
resource "azurerm_virtual_network" "vnet" {  
  name            = "${var.prefix}-vnet"  
  location        = azurerm_resource_group.hashitraining.location  
  address_space   = [var.address_space]  
  resource_group_name = azurerm_resource_group.hashitraining.name  
}
```

Variables & Outputs

Terraform

Variables

- Variables allow to bring more dynamic behavior to your Terraform code
- Defined on your code work as variables and parameters
- Variables can have default settings
- If you omit the default, the user will be prompted to enter a value

Variables

- Variables must have a **"name"** property
- Variables can have a **"description"** and **"default"** properties

```
variable "prefix" {  
    description = "This prefix will be included in the name of most resou  
}  
  
variable "location" {  
    description = "The region where the virtual network is created."  
    default     = "centralus"  
}
```

How to use variables?

- You may use variables value in two ways, direct or using interpolation
- Direct mode, you use the name of the variable with '**var.**' as prefix

```
resource "azurerm_virtual_network" "vnet" {  
  address_space = var.address_space  
  location      = var.location  
  name          = var.name  
  resource_group_name = var.resource_group  
  dns_servers    = var.dns_servers  
}
```

How to use variables?

- Using interpolation, you need to put your variable inside '`${ ... }`' and this code needs to be inside double quotes

```
resource "azurerm_virtual_network" "vnet" {  
  address_space = var.address_space  
  location      = var.location  
  name          = "${var.name}-vnet"  
  resource_group_name = "${var.resource_group}-rg"  
  dns_servers     = var.dns_servers  
}
```

How to set variables value?

- You can set variables values following the next precedence list
 1. Command line flag
 2. Configuration file
 3. Environment variable
 4. Default Config
 5. User manual entry - if not specified, prompt the user for entry

Outputs

- With outputs you can show or pass data that you gather after resources creation
- These values can be used as inputs for another resources to be used as properties values
- The outputs are automatically shown as log on the standard output

```
output "Vault_Server_URL" {  
    value = "http://${azurerm_public_ip.vault-pip.fqdn}:8200"  
}  
  
output "MySQL_Server_FQDN" {  
    value = azurerm_mysql_server.mysql.fqdn  
}  
  
output "catapp_url" {  
    value = "http://${azurerm_public_ip.catapp-pip.fqdn}"  
}
```


Outputs: Sensitive information

- You may protect your outputs to not be automatically write on the standard output marking as sensitive data
- The providers already do that on their resources so if you want to output a password/token/etc. you will see a redacted string
- To mark a output as sensitive, add the property '**sensitive = true**' on output definition

```
output "name" {  
  description = "The name of the newly created vNet"  
  value       = azurerm_virtual_network.vnet.name  
  sensitive   = true  
}
```

Organize your code

Terraform

Organize your code

- In Terraform you don't have mandatory way to organize your code
- Terraform CLI takes all input files and internally merge all the code
- But there is a recommended way to structure your code and is followed by all community

Best practices to author code

- **main.tf**: Main structure in order to deploy the complete infrastructure directly or through calls to modules
- **output.tf**: Identify each return value of a Terraform module
- **variables.tf**: Define variables required (or not) in main.tf, we need to define the var type and it allows to you assign a default value.
- **versions.tf**: Define minimum version required for terraform and providers in a module
- **terraform.tfvars**: Assign values to variables in a file instead of arguments in a command



Terraform lock file

- After you execute terraform init command, all providers/provisioners are downloaded to .terraform folder
- Additionally, a terraform lock file (named XXXXX) is created with hash and version information about all downloaded content
- This file should be treated as source code and should be pushed to your repo to keep track of providers versions

Terraform State

Terraform

Terraform State

- **Terraform State** is a critical concept in Terraform
- Represents a **snapshot of your infrastructure** at a specific point in time
- Terraform uses this state file to **keep track of resources** you've deployed to your providers
- With it ensures that it knows the current state of your infrastructure
- Terraform State is commonly named as **tfstate**

Terraform State: Why is important?

- When you use Terraform to create resources, it needs to remember:
 - Which resources it has already created
 - The properties of those resources (IDs, IP addresses, configuration details, etc.).
 - The relationship between your resources, like which instances are attached to which networks.
- Terraform **stores this information in a state file**, usually named `terraform.tfstate`
- This state file is used every time you run commands like **terraform plan** or **terraform apply**
- It's used to determine the **differences between your desired configuration and the actual resources** in your cloud environment

Terraform State: Terraform State format

- Terraform state file is a JSON file that must be only managed by Terraform
- You should never update this file manually
- Terraform (using Backend Providers) have different ways to make sure the tfstate was not changed manually

```
{
  "version": 4,
  "resources": [
    {
      "type": "aws_instance",
      "name": "my_server",
      "instances": [
        {
          "attributes": {
            "id": "i-0abcdef1234567890",
            "ami": "ami-12345678",
            "instance_type": "t2.micro",
            "public_ip": "54.12.34.56"
          }
        }
      ]
    }
  ]
}
```

Terraform State: How it works?

- **Initial State:** When you run **terraform apply** for the first time, Terraform creates the resources and writes their current state to **terraform.tfstate**.
- **State Updates:** On subsequent runs, Terraform compares the existing **terraform.tfstate** with your updated configuration files to determine what changes are needed (if any).
- **Tracking Changes:** Terraform uses this state to efficiently plan and apply only the necessary changes, minimizing disruption to your infrastructure.

Terraform State: Plan output

- When you run `terraform plan` and Terraform makes this comparison between your source code, terraform state and real resources you get a list of operations to be done

```
+   create
-   destroy
-/+  replace
~   update in-place
```

Demo – Your first Terraform Code

Terraform

Lab 01 – Your first Terraform Code

Terraform

